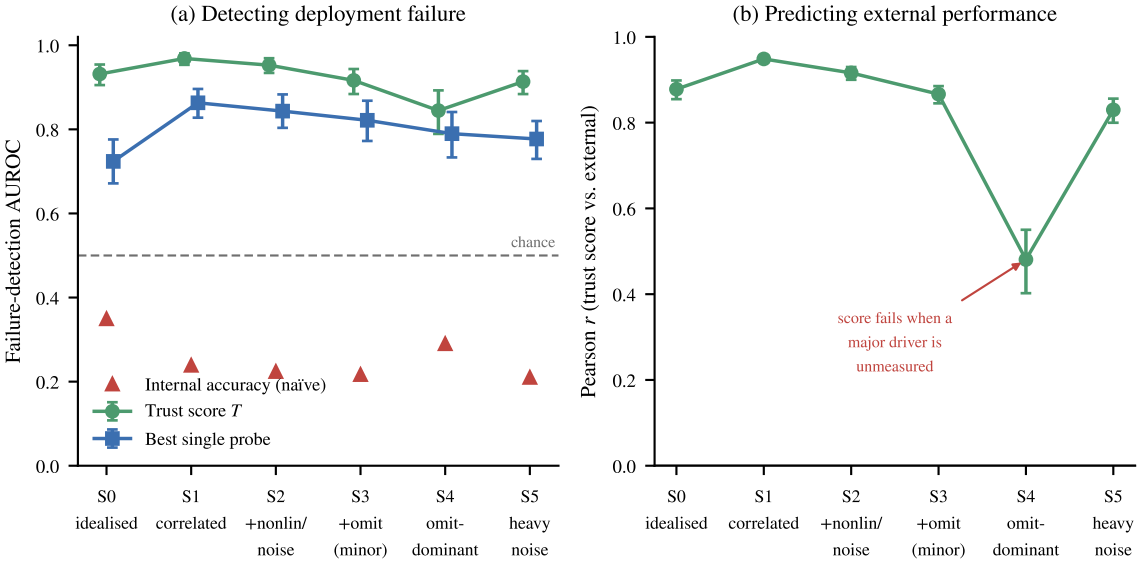
(a) Detecting deployment failure
Failure-detection AUROC
1.0
0.8
0.6
0.4
0.2
0.0
chance
Internal accuracy (naïve)
Trust score T
Best single probe
S0 idealised
S1 correlated
S2 +nonlin/ noise
S3 +omit (minor)
S4 omit-dominant
S5 heavy noise
(b) Predicting external performance
Pearson r (trust score vs. external)
score fails when a major driver is unmeasured